

NSD NOSQL DAY02

- 1. [案例1：部署redis集群](#)
- 2. [案例2：管理redis集群](#)

1 案例1：部署redis集群

1.1 问题

- 具体要求如下：
 - 准备集群环境
 - 安装redis并创建集群
 - 查看集群信息
 -
- 1.

1.2 方案

搭建redis集群，拓扑规划如图-1所示：

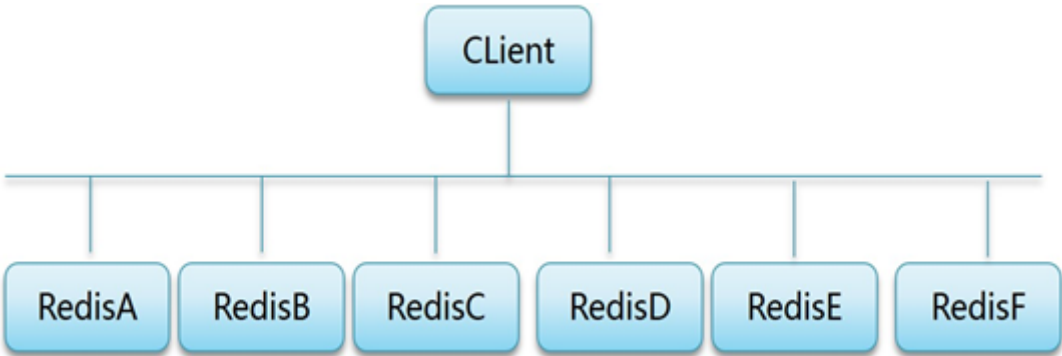


图 - 1

IP，端口规划如表-1所示：

表-1

主机名	IP 地址	端口号
redisA	192.168.4.51	6351
redisB	192.168.4.52	6352
redisC	192.168.4.53	6353
redisD	192.168.4.54	6354
redisE	192.168.4.55	6355
redisF	192.168.4.56	6356

1.3 步骤

实现此案例需要按照如下步骤进行。

[Top](#)

步骤一：准备集群环境

1) 按照表-1配置主机名, ip地址, 配置yum源(系统源) 这里不再操作

2) 把redis的软件包传到6台数据库服务器上面, 安装redis服务器, 六台服务器同样操作(以51为例)

```

01. [ root@redisA ~] # yum -y install gcc gcc-c++ make
02. [ root@redisA ~] # cd redis
03. redis/      redis-cluster/
04. [ root@redisA ~] # cd redis/
05. [ root@redisA redis] # ls
06. Inmp redis-4.0.8.tar.gz
07. [ root@redisA redis] # tar -xf redis-4.0.8.tar.gz
08. [ root@redisA redis] # cd redis-4.0.8/
09. [ root@redisA redis-4.0.8] # make && make install
10. [ root@redisA redis-4.0.8] # ./utils/install_server.sh
11. Welcome to the redis service installer
12. This script will help you easily set up a running redis server
13.
14. Please select the redis port for this instance: [ 6379]
15. Selecting default: 6379
16. Please select the redis config file name [ /etc/redis/6379.conf ]
17. Selected default - /etc/redis/6379.conf
18. Please select the redis log file name [ /var/log/redis_6379.log]
19. Selected default - /var/log/redis_6379.log
20. Please select the data directory for this instance [ /var/lib/redis/6379]
21. Selected default - /var/lib/redis/6379
22. Please select the redis executable path [ /usr/local/bin/redis-server]
23. Selected config:
24. Port      : 6379
25. Config file : /etc/redis/6379.conf
26. Log file   : /var/log/redis_6379.log
27. Data dir    : /var/lib/redis/6379
28. Executable  : /usr/local/bin/redis-server
29. Cli Executable : /usr/local/bin/redis-cli
30. Is this ok? Then press ENTER to go on or Ctrl-C to abort.
31. Copied /tmp/6379.conf => /etc/init.d/redis_6379
32. Installing service...
33.
34.
35. Successfully added to chkconfig!
36. Successfully added to runlevels 345!
37. Starting Redis server...

```

[Top](#)

```

38.  Installation successful! //安装成功
39.
40.  [root@redisA redis-4.0.8] # ss - antlp | grep 6379 //查看时有端口
41.  LISTEN 0 128 127.0.0.1:6379 *: * users: ( ( "redis-server"

```

2) 修改配置文件，6台redis服务器都要修改（以51为例子）

```

01.  [root@redisA redis-4.0.8] # /etc/init.d/redis_6379 stop
02.  //停止已经开启的redis服务
03.  Stopping ...
04.  Waiting for Redis to shutdown ...
05.  Redis stopped
06.
07.
08.  [root@redisA redis-4.0.8] # vim /etc/redis/6379.conf
09.  ...
10.  bind 192.168.4.51 //修改ip
11.  port 6351 //不允许相同，只指定物理接口的ip
12.  daemonize yes //以守护进程方式运行
13.  pidfile /var/run/redis_6351.pid
14.  cluster-enabled yes //是否启用集群，前提是以守护进程方式运行
15.  cluster-config-file nodes-6351.conf
16.  //存储集群信息的配置文件，自动生成，不允许相同
17.  cluster-node-timeout 5000 //集群节点通信超时时间
18.  ...
19.  [root@redisA redis-4.0.8] # /etc/init.d/redis_6379 start //启动服务
20.  Starting Redis server...
21.  [root@redisA redis-4.0.8] # ss - antlp | grep 6351 //查看有端口
22.  LISTEN 0 128 192.168.4.51:6351 *: * users: ( ( "redis-ser
23.  LISTEN 0 128 192.168.4.51:16351 *: * users: ( ( "redis-ser
24.  [root@redisA redis-4.0.8] # ps - C redis
25.  PID TTY TIME CMD

```

注意：其他几台主机在修改时请注意ip，端口等的修改，不要和51主机的一样

3) 关闭防火墙51-56主机（以51为例子）

[Top](#)

```

01.  [root@redisA redis-4.0.8] # getenforce

```

02. Permissive
- 03.
04. [root@redisA redis- 4.0.8] # systemctl disable firewalld
05. //关闭防火墙不自启

4) 查看集群信息

01. [root@redisA redis- 4.0.8] # redis- cli - h 192.168.4.51 - p 6351
02. 192.168.4.51: 6351> ping
03. PONG
04. 192.168.4.51: 6351> cluster info
05. cluster_state: fail
06. cluster_slots_assigned: 0
07. cluster_slots_ok: 0
08. cluster_slots_pfail: 0
09. cluster_slots_fail: 0
10. cluster_known_nodes: 1
11. cluster_size: 0
12. ...
- 13.
14. 192.168.4.51: 6351> cluster nodes
15. f81f997d5ed988ec1587558e78d5f7dbc96abcbf : 6351@16351 my self , master - 0 0 0 connec

步骤二：创建集群（在任意一台上执行创建集群的脚本都可以）这里在51上面执行

1) 部署ruby脚本运行环境（在51上面执行）

01. [root@redisA redis- 4.0.8] # cd /root/redis- cluster/
02. [root@redisA redis- cluster] # ls
03. redis- 3.2.1.gem ruby- devel- 2.0.0.648- 30.el7.x86_64.rpm
04. [root@redisA redis- cluster] # yum - y install ruby rubygems
- 05.
06. [root@redisA redis- cluster] # rpm - ivh -- nodeps \
07. ruby- devel- 2.0.0.648- 30.el7.x86_64.rpm
08. warning: ruby- devel- 2.0.0.648- 30.el7.x86_64.rpm: Header V3 RSA / SHA 256 Signature, ke
09. Preparing... ##### [100%]
10. Updating / installing... [Top](#)
11. 1: ruby- devel- 2.0.0.648- 30.el7 ##### [100%]
12. [root@redisA redis- cluster] # which gem

13. `/usr/bin/gem`
14. `[root@redisA redis-cluster] # gem install redis`
15. Successfully installed redis- 3.2.1
16. Parsing documentation for redis- 3.2.1
17. Installing ri documentation for redis- 3.2.1
18. 1 gem installed

2) 生成创建集群的脚本

01. `[root@redisA redis-cluster] # cd /root/redis/redis- 4.0.8/src/`
02. `[root@redisA src] # cp redis-trib.rb /usr/local/bin/`
03. `[root@redisA src] # ll /usr/local/bin/redis-trib.rb`
04. `-rw-r--r-x. 1 root root 65991 Sep 27 16:12 /usr/local/bin/redis-trib.rb`

3) 创建集群

01. `[root@redisA src] # redis-trib.rb create --replicas 1 \`
02. `192.168.4.51:6351 192.168.4.52:6352 \`
03. `192.168.4.53:6353 192.168.4.54:6354 \`
04. `192.168.4.55:6355 192.168.4.56:6356`
05. `// --replicas 1 给每一个主配置一个从库`
06. `[root@redisA log] # redis-trib.rb create --replicas 1 192.168.4.51:6351 192.168.4.52:6352`
07. `>>> Creating cluster`
08. `>>> Performing hash slots allocation on 6 nodes...`
09. `Using 3 masters:`
10. `192.168.4.51:6351`
11. `192.168.4.52:6352`
12. `192.168.4.53:6353`
13. `Adding replica 192.168.4.55:6355 to 192.168.4.51:6351`
14. `Adding replica 192.168.4.56:6356 to 192.168.4.52:6352`
15. `Adding replica 192.168.4.54:6354 to 192.168.4.53:6353`
16. `...`
17. `...`
18. `...`
19. `[OK] All nodes agree about slots configuration.`
20. `>>> Check for open slots...`
21. `>>> Check slots coverage...`
22. `[OK] All 16384 slots covered.`

[Top](#)

4) 查看集群信息，任意一台主机访问本机的redis服务查看即可

cluster info 查看集群信息

cluster nodes 查看集群节点信息

```

01. [root@redisA log] # redis-cli -h 192.168.4.52 -p 6352
02. 192.168.4.52:6352> CLUSTER INFO
03. cluster_state:ok //状态
04. cluster_slots_assigned:16384
05. cluster_slots_ok:16384
06. cluster_slots_pfail:0
07. cluster_slots_fail:0
08. cluster_known_nodes:6
09. cluster_size:3
10. cluster_current_epoch:6
11. cluster_my_epoch:2
12. cluster_stats_messages_ping_sent:367
13. cluster_stats_messages_pong_sent:327
14. cluster_stats_messages_meet_sent:5
15. cluster_stats_messages_sent:699
16. cluster_stats_messages_ping_received:327
17. cluster_stats_messages_pong_received:372
18. cluster_stats_messages_received:699
19.
20. 192.168.4.52:6352> CLUSTER NODES //查看集群节点信息
21. 63afbb5e7d63b1f142358634578a3488e3c9e634 192.168.4.54:6354@16354 slave bc5c4e082a5a3391b634cf433a6486c867cf c44b 192.168.4.53:6353@16353 master - 0 15380
22. bc5c4e082a5a3391b634cf433a6486c867cf c44b 192.168.4.53:6353@16353 master - 0 15380
23. 28e06c5f24a2b6c6412f81369e09bc9653cc51ff 192.168.4.56:6356@16356 slave 8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 myself ,master -
24. 7e8d9121f44d8331ff55b45c218b87df9bda1b70 192.168.4.55:6355@16355 slave a3083123bc5c87a76aab2655171634d4ee84f418 192.168.4.51:6351@16351 master - 0 15380
25. 8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 myself ,master -
26. a3083123bc5c87a76aab2655171634d4ee84f418 192.168.4.51:6351@16351 master - 0 15380
27. 192.168.4.52:6352>

```

5) 测试集群

命令：

redis-cli -c -h ip地址 -p 端口

[Top](#)

```

01. [root@redisA log] # redis-cli -c -h 192.168.4.51 -p 6351

```

```

02. 192.168.4.51:6351> set name jim
03. - > Redirected to slot [ 5798] located at 192.168.4.52:6352
04. OK
05. 192.168.4.52:6352> get name
06. "jim"
07. 192.168.4.52:6352> set class linux
08. OK
09. 192.168.4.52:6352> get class
10. "linux"
11. 192.168.4.52:6352> set pay 26800
12. - > Redirected to slot [ 4013] located at 192.168.4.51:6351
13. OK
14. 192.168.4.51:6351> get pay
15. "26800"

```

集群不能用的情况：

有半数或者半数以上的主库机器挂掉，集群就不能用了

把一个从库升级成主，没有从库，集群不能用（前提是：有半数或者半数以上的主库机器挂掉）

一个主库挂掉，它的从库自动顶替为主库，正常使用（前提是：有半数或者半数以上的主库机器能用），挂掉的主库修复好后，会成为从库，不会抢占为主

6）集群节点选举策略（三主，三从）

停止某个主库的redis服务，对应的从库会自动升级为主库

先查看节点信息的主从情况

```

01. [ root@redisA log] # redis-cli -c -h 192.168.4.51 -p 6351
02. 192.168.4.51:6351> CLUSTER nodes
03. ...
04. 8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 master - 0 1538x
05. 28e06c5f24a2b6c6412f81369e09bc9653cc51f 192.168.4.56:6356@16356 slave 8568fbd73c
06. ...
07. 192.168.4.51:6351>

```

看谁是谁的从库，如：

看节点前后的编号id是否有相同的

如：8568fbd73cb296cad6915d524e34761b2114af47

发现52的从库为56

停止主库52

[Top](#)

```

01. [root@redisA log] # redis-cli -h 192.168.4.52 -p 6352 shutdown
02. [root@redisA log] # redis-cli -c -h 192.168.4.51 -p 6351
03. 192.168.4.51:6351> CLUSTER NODES
04. ...
05. 8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 master,fail - 153
06. 28e06c5f24a2b6c6412f81369e09bc9653cc51f 192.168.4.56:6356@16356 master - 0 15380
07. ...

```

开启52，发现52成为从库

```

01. [root@redisB redis-4.0.8] # /etc/init.d/redis_6352 start
02. Starting Redis server...
03. [root@redisA log] # redis-cli -c -h 192.168.4.51 -p 6351
04. 192.168.4.51:6351> CLUSTER NODES
05. 8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 slave 28e06c5f2

```

2 案例2：管理redis集群

2.1 问题

- 具体要求如下：
- 练习添加主机
- 练习删除主机

2.2 步骤

实现此案例需要按照如下步骤进行。

步骤一：添加主机

1) 部署一台新redis服务器，ip为192.168.4.58，装包，初始化，启用集群配置，重启服务（这里之前已经操作过，不会的可以参考案例1）

2) 添加集群4.58（添加master节点）

格式：redis-trib.rb 选项 参数

选项：add-node 添加主机（不指定角色为主）

由于之前是在51上面创建ruby脚本，所以只有51上面有redis-trib.rb命令，在51上面执行

```

01. [root@redisA ~] # redis-trib.rb add-node 192.168.4.58:6358 192.168.4.51:6351
02. >>> Adding node 192.168.4.58:6358 to cluster 192.168.4.51:6351 Top
03. >>> Performing Cluster Check ( using node 192.168.4.51:6351)
04. S: a3083123bc5c87a76aab2655171634d4ee84f418 192.168.4.51:6351

```



```

05. slots: ( 0 slots) slave
06. replicates 7e8d9121f 44d8331ff 55b45c218b87df 9bda1b70
07. M: 7e8d9121f 44d8331ff 55b45c218b87df 9bda1b70 192.168.4.55: 6355
08. slots: 0- 5460 ( 5461 slots) master
09. 1 additional replica( s)
10. S: 8568f bd73cb296cad6915d524e34761b2114af 47 192.168.4.52: 6352
11. slots: ( 0 slots) slave
12. replicates 28e06c5f 24a2b6c6412f 81369e09bc9653cc51ff
13. M: bc5c4e082a5a3391b634cf 433a6486c867cf c44b 192.168.4.53: 6353
14. slots: 10923- 16383 ( 5461 slots) master
15. 1 additional replica( s)
16. S: 63afb b5e7d63b1f 142358634578a3488e3c9e634 192.168.4.54: 6354
17. slots: ( 0 slots) slave
18. replicates bc5c4e082a5a3391b634cf 433a6486c867cf c44b
19. M: 28e06c5f 24a2b6c6412f 81369e09bc9653cc51ff 192.168.4.56: 6356
20. slots: 5461- 10922 ( 5462 slots) master
21. 1 additional replica( s)
22. [ OK] All nodes agree about slots configuration.
23. >>> Check for open slots...
24. >>> Check slots coverage...
25. [ OK] All 16384 slots covered.
26. >>> Send CLUSTER MEET to node 192.168.4.58: 6358 to make it join the cluster.
27. [ OK] New node added correctly.

```

3) 检查集群主机的状态信息

选项 : check 检查权限

```

01. [ root@redisA ~] # redis-trib.rb check 192.168.4.58: 6358 //查看状态
02. >>> Performing Cluster Check ( using node 192.168.4.58: 6358)
03. M: c5e0da48f 335c46a2ec199f aa99b830f 537dd8a0 192.168.4.58: 6358
04. slots: ( 0 slots) master //发现没有hash槽
05. 0 additional replica( s)
06. M: 7e8d9121f 44d8331ff 55b45c218b87df 9bda1b70 192.168.4.55: 6355
07. slots: 0- 5460 ( 5461 slots) master
08. 1 additional replica( s)
09. ...
10. S: a3083123bc5c87a76aab2655171634d4ee84f 418 192.168.4.51: 6351
11. slots: ( 0 slots) slave
12. replicates 7e8d9121f 44d8331ff 55b45c218b87df 9bda1b70
13. [ OK] All nodes agree about slots configuration.

```

[Top](#)

14. >>> Check for open slots...
15. >>> Check slots coverage...
16. [OK] All 16384 slots covered.

4) 手动对集群进行分片迁移

选项：reshard 重新分配hash槽

01. [root@redisA ~] # redis-trib.rb reshard 192.168.4.58:6358
02. How many slots do you want to move (from 1 to 16384) ?4096
03. //拿出多少个hash 槽给主机192.168.4.58
04. What is the receiving node ID? c5e0da48f335c46a2ec199faa99b830f537dd8a0
05. //主机192.168.4.58的id值
06. Source node #1: all //从当前所有的主里面获取hash槽
07. Do you want to proceed with the proposed reshard plan (yes/no) ?yes
08. ...
09. Moving slot 12283 from 192.168.4.53:6353 to 192.168.4.58:6358:
10. Moving slot 12284 from 192.168.4.53:6353 to 192.168.4.58:6358:
11. Moving slot 12285 from 192.168.4.53:6353 to 192.168.4.58:6358:
12. Moving slot 12286 from 192.168.4.53:6353 to 192.168.4.58:6358:
13. Moving slot 12287 from 192.168.4.53:6353 to 192.168.4.58:6358:

再次查看发现4.58有4096个hash slot

01. [root@redisA ~] # redis-trib.rb check 192.168.4.58:6358
02. >>> Performing Cluster Check (using node 192.168.4.58:6358)
03. M: c5e0da48f335c46a2ec199faa99b830f537dd8a0 192.168.4.58:6358
04. slots: 0-1364,5461-6826,10923-12287 (4096 slots) master
05. 0 additional replica(s)

5) 删除master角色的主机

先删除主机占用的hash槽

01. [root@redisA ~] # redis-trib.rb reshard 192.168.4.58:6358
02. How many slots do you want to move (from 1 to 16384) ?4096
03. //移除hash 槽的个数
04. What is the receiving node ID? bc5c4e082a5a3391b634cf433a6486c867cf c44b
05. //要移动给谁的id即目标主机 (这里可以随机写一个master的ID)

[Top](#)

```

06. Source node #1: c5e0da48f335c46a2ec199f aa99b830f537dd8a0
07. //从谁那移动即源主机 (这里写4.58的ID)
08. Source node #2: done //设置完毕
09. ...
10. Moving slot 12282 from c5e0da48f335c46a2ec199f aa99b830f537dd8a0
11. Moving slot 12283 from c5e0da48f335c46a2ec199f aa99b830f537dd8a0
12. Moving slot 12284 from c5e0da48f335c46a2ec199f aa99b830f537dd8a0
13. Moving slot 12285 from c5e0da48f335c46a2ec199f aa99b830f537dd8a0
14. Moving slot 12286 from c5e0da48f335c46a2ec199f aa99b830f537dd8a0
15. Moving slot 12287 from c5e0da48f335c46a2ec199f aa99b830f537dd8a0
16. Do you want to proceed with the proposed reshard plan ( yes/no )? yes //提交
17. ...
18. Moving slot 12282 from 192.168.4.58:6358 to 192.168.4.53:6353:
19. Moving slot 12283 from 192.168.4.58:6358 to 192.168.4.53:6353:
20. Moving slot 12284 from 192.168.4.58:6358 to 192.168.4.53:6353:
21. Moving slot 12285 from 192.168.4.58:6358 to 192.168.4.53:6353:
22. Moving slot 12286 from 192.168.4.58:6358 to 192.168.4.53:6353:
23. Moving slot 12287 from 192.168.4.58:6358 to 192.168.4.53:6353:

```

删除集群主机4.58(删除之后redis服务自动关闭)

```

01. [ root@redisA ~] # redis-trib.rb del-node 192.168.4.58:6358 \
02. c5e0da48f335c46a2ec199f aa99b830f537dd8a0 //删除谁+删除的id
03. >>> Removing node c5e0da48f335c46a2ec199f aa99b830f537dd8a0 from cluster 192.168.4.
04. >>> Sending CLUSTER FORGET messages to the cluster...
05. >>> SHUTDOWN the node.

```

6) 添加从节点主机，随机添加

```

01. [ root@redisA ~] # redis-trib.rb add-node --slave \
02. 192.168.4.57:6357 192.168.4.51:6351
03. >>> Adding node 192.168.4.57:6357 to cluster 192.168.4.51:6351
04. >>> Performing Cluster Check ( using node 192.168.4.51:6351)
05. .....
06. .....
07. [ OK ] All 16384 slots covered.
08. Automatically selected master 192.168.4.51:6351
09. >>> Send CLUSTER MEET to node 192.168.4.57:6357 to make it join the cluster.

```

[Top](#)

10. Waiting for the cluster to join.
11. >>> Configure node as replica of 192.168.4.51:6351
12. [OK] New node added correctly .

7) 移除从节点，从节点主机没有槽位范围，直接移除即可

命令格式：

redis-trib.rb del-node 192.168.4.57:6357 主机id值

01. [root@redisA ~] # redis-trib.rb del-node 192.168.4.57:6357 \
02. f6649ea99b2f01faca26217691222c17a3854381
03. >>> Removing node f6649ea99b2f01faca26217691222c17a3854381
04. from cluster 192.168.4.57:6351
05. >>> Sending CLUSTER FORGET messages to the cluster...
06. >>> SHUT DOWN the node.

[Top](#)